

act2_uresti

September 11, 2025

1 Actividad Modulo 2

1.0.1 Segunda entrega

Yose Miguel Sotomayor Carneado A0150908

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.datasets import fetch_kddcup99
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, ConfusionMatrixDisplay
from sklearn.preprocessing import StandardScaler, LabelEncoder, PowerTransformer

from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression

import warnings
warnings.filterwarnings("ignore")
```

```
[3]: data = fetch_kddcup99(as_frame=True)

X = data.data
y = data.target.astype(str)
```

```
[4]: encoder = LabelEncoder()

y = encoder.fit_transform(y)
np.unique(y)
```

```
[4]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
          17, 18, 19, 20, 21, 22])
```

```
[60]: pd.DataFrame(y).value_counts()
```

```
[60]: 0
      18    280790
      9    107201
     11     97278
      0     2203
     17     1589
      5     1247
     15     1040
     21     1020
     20      979
     14      264
     10      231
      3       53
      1       30
      6       21
     22       20
      4       12
     16       10
      7        9
      2        8
      8        7
     13        4
     12        3
     19        2
      Name: count, dtype: int64
```

```
[5]: X
```

```
[5]:      duration protocol_type  service  flag  src_bytes  dst_bytes  land  \
0           0      b'tcp'  b'http'  b'SF'      181      5450    0
1           0      b'tcp'  b'http'  b'SF'      239       486    0
2           0      b'tcp'  b'http'  b'SF'      235     1337    0
3           0      b'tcp'  b'http'  b'SF'      219     1337    0
4           0      b'tcp'  b'http'  b'SF'      217     2032    0
...      ...      ...      ...      ...      ...      ...
494016      0      b'tcp'  b'http'  b'SF'      310     1881    0
494017      0      b'tcp'  b'http'  b'SF'      282     2286    0
494018      0      b'tcp'  b'http'  b'SF'      203     1200    0
494019      0      b'tcp'  b'http'  b'SF'      291     1200    0
494020      0      b'tcp'  b'http'  b'SF'      219     1234    0
```

```
wrong_fragment  urgent  hot  ...  dst_host_count  dst_host_srv_count  \
```

0	0	0	0	...	9	9
1	0	0	0	...	19	19
2	0	0	0	...	29	29
3	0	0	0	...	39	39
4	0	0	0	...	49	49
...
494016	0	0	0	...	86	255
494017	0	0	0	...	6	255
494018	0	0	0	...	16	255
494019	0	0	0	...	26	255
494020	0	0	0	...	6	255

	dst_host_same_srv_rate	dst_host_diff_srv_rate	\
0	1.0	0.0	
1	1.0	0.0	
2	1.0	0.0	
3	1.0	0.0	
4	1.0	0.0	
...	
494016	1.0	0.0	
494017	1.0	0.0	
494018	1.0	0.0	
494019	1.0	0.0	
494020	1.0	0.0	

	dst_host_same_src_port_rate	dst_host_srv_diff_host_rate	\
0	0.11	0.0	
1	0.05	0.0	
2	0.03	0.0	
3	0.03	0.0	
4	0.02	0.0	
...	
494016	0.01	0.05	
494017	0.17	0.05	
494018	0.06	0.05	
494019	0.04	0.05	
494020	0.17	0.05	

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	
...	
494016	0.0	0.01	0.0	
494017	0.0	0.01	0.0	

494018	0.06	0.01	0.0
494019	0.04	0.01	0.0
494020	0.0	0.01	0.0

	dst_host_srv_rerror_rate
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
494016	0.0
494017	0.0
494018	0.0
494019	0.0
494020	0.0

[494021 rows x 41 columns]

[10]: X.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 494021 entries, 0 to 494020
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   duration                             494021 non-null  object
1   protocol_type                         494021 non-null  object
2   service                              494021 non-null  object
3   flag                                  494021 non-null  object
4   src_bytes                            494021 non-null  object
5   dst_bytes                            494021 non-null  object
6   land                                  494021 non-null  object
7   wrong_fragment                       494021 non-null  object
8   urgent                               494021 non-null  object
9   hot                                  494021 non-null  object
10  num_failed_logins                    494021 non-null  object
11  logged_in                           494021 non-null  object
12  num_compromised                      494021 non-null  object
13  root_shell                           494021 non-null  object
14  su_attempted                        494021 non-null  object
15  num_root                             494021 non-null  object
16  num_file_creations                  494021 non-null  object
17  num_shells                           494021 non-null  object
18  num_access_files                    494021 non-null  object
19  num_outbound_cmds                   494021 non-null  object
20  is_host_login                       494021 non-null  object
21  is_guest_login                       494021 non-null  object
```

22	count	494021	non-null	object
23	srv_count	494021	non-null	object
24	serror_rate	494021	non-null	object
25	srv_serror_rate	494021	non-null	object
26	rerror_rate	494021	non-null	object
27	srv_rerror_rate	494021	non-null	object
28	same_srv_rate	494021	non-null	object
29	diff_srv_rate	494021	non-null	object
30	srv_diff_host_rate	494021	non-null	object
31	dst_host_count	494021	non-null	object
32	dst_host_srv_count	494021	non-null	object
33	dst_host_same_srv_rate	494021	non-null	object
34	dst_host_diff_srv_rate	494021	non-null	object
35	dst_host_same_src_port_rate	494021	non-null	object
36	dst_host_srv_diff_host_rate	494021	non-null	object
37	dst_host_serror_rate	494021	non-null	object
38	dst_host_srv_serror_rate	494021	non-null	object
39	dst_host_rerror_rate	494021	non-null	object
40	dst_host_srv_rerror_rate	494021	non-null	object

dtypes: object(41)
memory usage: 154.5+ MB

```
[11]: X.isna().sum()
```

```
[11]: duration          0
      protocol_type     0
      service           0
      flag              0
      src_bytes         0
      dst_bytes         0
      land              0
      wrong_fragment    0
      urgent            0
      hot               0
      num_failed_logins 0
      logged_in         0
      num_compromised   0
      root_shell        0
      su_attempted      0
      num_root          0
      num_file_creations 0
      num_shells         0
      num_access_files  0
      num_outbound_cmds 0
      is_host_login     0
      is_guest_login    0
      count             0
```

```

srv_count          0
serror_rate        0
srv_serror_rate     0
rerror_rate        0
srv_rerror_rate     0
same_srv_rate       0
diff_srv_rate       0
srv_diff_host_rate  0
dst_host_count      0
dst_host_srv_count  0
dst_host_same_srv_rate 0
dst_host_diff_srv_rate 0
dst_host_same_src_port_rate 0
dst_host_srv_diff_host_rate 0
dst_host_serror_rate 0
dst_host_srv_serror_rate 0
dst_host_rerror_rate 0
dst_host_srv_rerror_rate 0
dtype: int64

```

```

[19]: # MAKE NUM CATEGORICAL and astype float

X['protocol_type'] = encoder.fit_transform(X['protocol_type'])
X['service'] = encoder.fit_transform(X['service'])
X['flag'] = encoder.fit_transform(X['flag'])

def make_float(x):
    return float(x)

for col in X.columns:
    X[col] = X[col].apply(make_float)

X

```

```

[19]:
   duration  protocol_type  service  flag  src_bytes  dst_bytes  land  \
0         0.0            1.0     22.0   9.0     181.0     5450.0   0.0
1         0.0            1.0     22.0   9.0     239.0      486.0   0.0
2         0.0            1.0     22.0   9.0     235.0     1337.0   0.0
3         0.0            1.0     22.0   9.0     219.0     1337.0   0.0
4         0.0            1.0     22.0   9.0     217.0     2032.0   0.0
...         ...           ...     ...   ...         ...     ...
494016     0.0            1.0     22.0   9.0     310.0     1881.0   0.0
494017     0.0            1.0     22.0   9.0     282.0     2286.0   0.0
494018     0.0            1.0     22.0   9.0     203.0     1200.0   0.0
494019     0.0            1.0     22.0   9.0     291.0     1200.0   0.0
494020     0.0            1.0     22.0   9.0     219.0     1234.0   0.0

```

	wrong_fragment	urgent	hot	...	dst_host_count	dst_host_srv_count	\
0	0.0	0.0	0.0	...	9.0	9.0	
1	0.0	0.0	0.0	...	19.0	19.0	
2	0.0	0.0	0.0	...	29.0	29.0	
3	0.0	0.0	0.0	...	39.0	39.0	
4	0.0	0.0	0.0	...	49.0	49.0	
...	
494016	0.0	0.0	0.0	...	86.0	255.0	
494017	0.0	0.0	0.0	...	6.0	255.0	
494018	0.0	0.0	0.0	...	16.0	255.0	
494019	0.0	0.0	0.0	...	26.0	255.0	
494020	0.0	0.0	0.0	...	6.0	255.0	

	dst_host_same_srv_rate	dst_host_diff_srv_rate	\
0	1.0	0.0	
1	1.0	0.0	
2	1.0	0.0	
3	1.0	0.0	
4	1.0	0.0	
...	
494016	1.0	0.0	
494017	1.0	0.0	
494018	1.0	0.0	
494019	1.0	0.0	
494020	1.0	0.0	

	dst_host_same_src_port_rate	dst_host_srv_diff_host_rate	\
0	0.11	0.00	
1	0.05	0.00	
2	0.03	0.00	
3	0.03	0.00	
4	0.02	0.00	
...	
494016	0.01	0.05	
494017	0.17	0.05	
494018	0.06	0.05	
494019	0.04	0.05	
494020	0.17	0.05	

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate	\
0	0.00	0.00	0.0	
1	0.00	0.00	0.0	
2	0.00	0.00	0.0	
3	0.00	0.00	0.0	
4	0.00	0.00	0.0	
...	
494016	0.00	0.01	0.0	

494017	0.00	0.01	0.0
494018	0.06	0.01	0.0
494019	0.04	0.01	0.0
494020	0.00	0.01	0.0

	dst_host_srv_error_rate
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
494016	0.0
494017	0.0
494018	0.0
494019	0.0
494020	0.0

[494021 rows x 41 columns]

```
[27]: X.hist(figsize=(25, 20))
plt.show()
```




```
[20]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

Como es un problema de clasificacion multiclase vamos a usar onehotencoder para las clases de la y

```
[54]: # Columnas categóricas
cat_cols = ['protocol_type', 'service', 'flag']

# Columnas numéricas
num_cols = [c for c in X.columns if c not in cat_cols]

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), num_cols)
    ]
)

transformer = Pipeline(
    steps=[
        ("power", PowerTransformer(method="yeo-johnson", standardize=True)),
    ]
)

clf = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('transform', transformer),
    ('classifier', LogisticRegression(max_iter=1000, penalty='l2',
↳ multi_class='auto' ))
])

clf.fit(X_train, y_train)
```

```
[54]: Pipeline(steps=[('preprocessor',
    ColumnTransformer(transformers=[('num', StandardScaler(),
    ['duration', 'src_bytes',
    'dst_bytes', 'land',
    'wrong_fragment', 'urgent',
    'hot', 'num_failed_logins',
    'logged_in',
    'num_compromised',
    'root_shell', 'su_attempted',
    'num_root',
```

```

        'num_file_creations',
        'num_shells',
        'num_access_files',
        'num_outbound_cmds',
        'is_host_login',
        'is_guest_login', 'count',
        'srv_count', 'serror_rate',
        'srv_serror_rate',
        'rerror_rate',
        'srv_rerror_rate',
        'same_srv_rate',
        'diff_srv_rate',
        'srv_diff_host_rate',
        'dst_host_count',
        'dst_host_srv_count',
...]])),
    ('transform', Pipeline(steps=[('power', PowerTransformer())])),
    ('classifier',
     LogisticRegression(max_iter=1000, multi_class='auto'))])

```

```
[55]: y_pred = clf.predict(X_test)
```

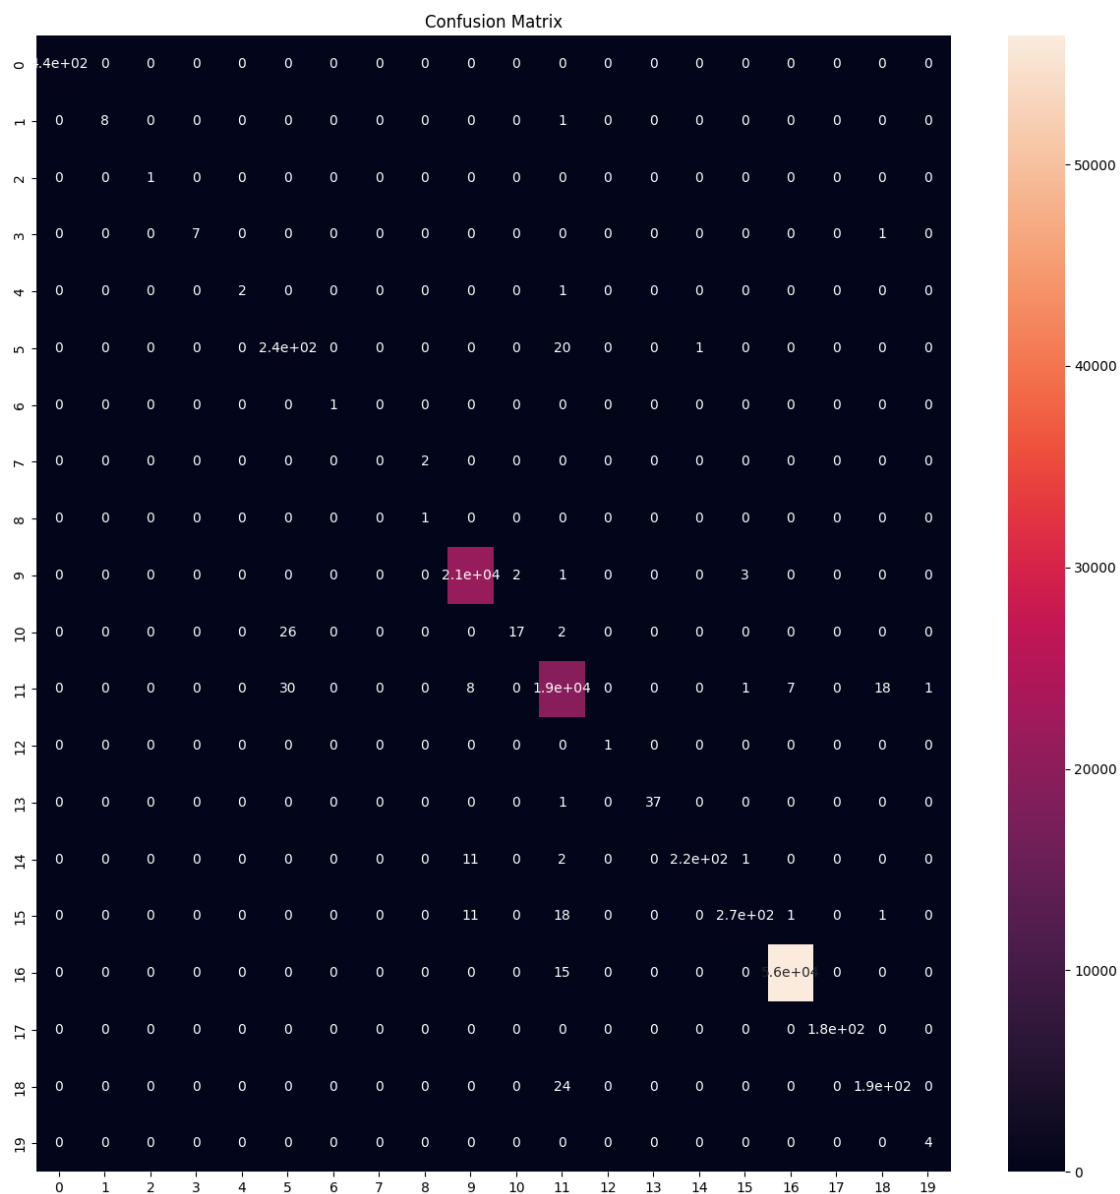
```
[56]: print(classification_report(y_pred, y_test))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	435
1	0.89	1.00	0.94	8
2	1.00	1.00	1.00	1
3	0.88	1.00	0.93	7
4	0.67	1.00	0.80	2
5	0.92	0.81	0.86	300
6	1.00	1.00	1.00	1
7	0.00	0.00	0.00	0
8	1.00	0.33	0.50	3
9	1.00	1.00	1.00	21318
10	0.38	0.89	0.53	19
11	1.00	1.00	1.00	19373
12	1.00	1.00	1.00	1
14	0.97	1.00	0.99	37
15	0.94	1.00	0.97	223
17	0.90	0.98	0.94	278
18	1.00	1.00	1.00	56395
20	1.00	1.00	1.00	185
21	0.89	0.91	0.90	214
22	1.00	0.80	0.89	5
accuracy			1.00	98805

macro avg	0.87	0.89	0.86	98805
weighted avg	1.00	1.00	1.00	98805

Podemos observar que los resultados son bastante buenos.

```
[57]: from sklearn.metrics import confusion_matrix
import seaborn as sns
plt.figure(figsize=(15, 15))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True)
plt.title("Confusion Matrix")
plt.show()
```



```
[58]: from sklearn.metrics import f1_score
print('F1 Macro' ,f1_score(y_test, y_pred, average='macro'))
print('F1 Micro' ,f1_score(y_test, y_pred, average='micro'))
print('F1 Weighted' ,f1_score(y_test, y_pred, average='weighted'))
```

F1 Macro 0.8621861403907598

F1 Micro 0.997874601487779

F1 Weighted 0.9978189895431658

Vemos que el modelo a grandes rasgos tiene unos buenos f1-score para evaluar el modelo con muchas clases y MUY mal balanceadas!